

CH -10

Unsupervised Learning and Clustering

By:

Arshad Farhad

20177716

Contents

- ◆ Supervised vs Unsupervised learning
 - ◆ Introduction to clustering
 - ◆ K-means Clustering
 - ◆ Hierarchical clustering
 - ◆ Conclusion
-

Supervised Vs Unsupervised Learning

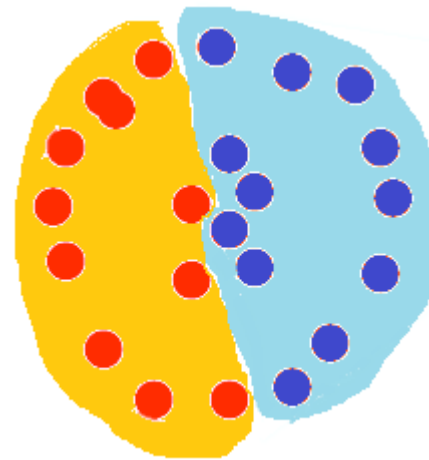
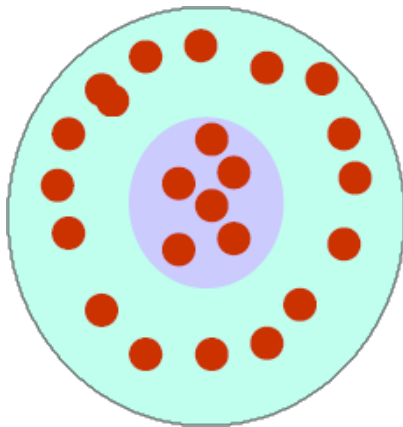
- ❑ **Supervised learning** is where you have input variables (x) and an output variable (Y) and you use an algorithm to learn the mapping function from the input to the output.

$$Y = f(X)$$

- ❑ *The goal* is to approximate the mapping function so well that when you have new input data (x) that you can predict the output variables (Y) for that data
 - ❑ **Unsupervised learning** is where you only have input data (X) and no corresponding output variables
 - ❑ *The goal* for unsupervised learning is to model the underlying structure or distribution in the data in order to learn more about the data.
 - ❑ Unsupervised learning problems can be further grouped into clustering and association problems.
 - **Clustering**
 - **Association**
-

What is clustering?

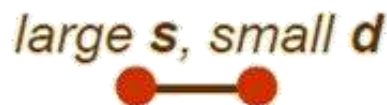
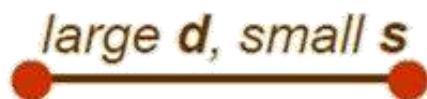
- The organization of unlabeled data into similarity groups called clusters.
- A cluster is a collection of data items which are “similar” between them, and “dissimilar” to data items in other clusters.



What do we need for clustering?

1. Proximity measure, *either*

- similarity measure $s(x_i, x_k)$: large if x_i, x_k are similar
- dissimilarity(or distance) measure $d(x_i, x_k)$: small if x_i, x_k are similar



2. Criterion function to evaluate a clustering



3. Algorithm to compute clustering

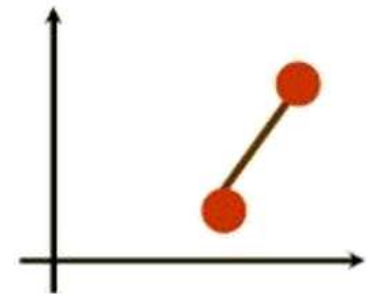
- For example, by optimizing the criterion function

Distance (dissimilarity) measures

- Euclidean distance between points i and j is the length of the line segment connecting them
- In Cartesian coordinates, if $\mathbf{i} = (i_1, i_2, \dots, i_n)$ and $\mathbf{q} = (q_1, q_2, \dots, q_n)$ then the distance (d) from i to j , or from j to i is given by:

- Euclidean distance

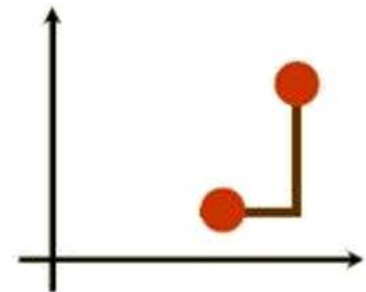
$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{\sum_{k=1}^d (\mathbf{x}_i^{(k)} - \mathbf{x}_j^{(k)})^2}$$



- Manhattan (city block) distance

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sum_{k=1}^d |\mathbf{x}_i^{(k)} - \mathbf{x}_j^{(k)}|$$

- approximation to Euclidean distance, cheaper to compute



Cluster Evaluation

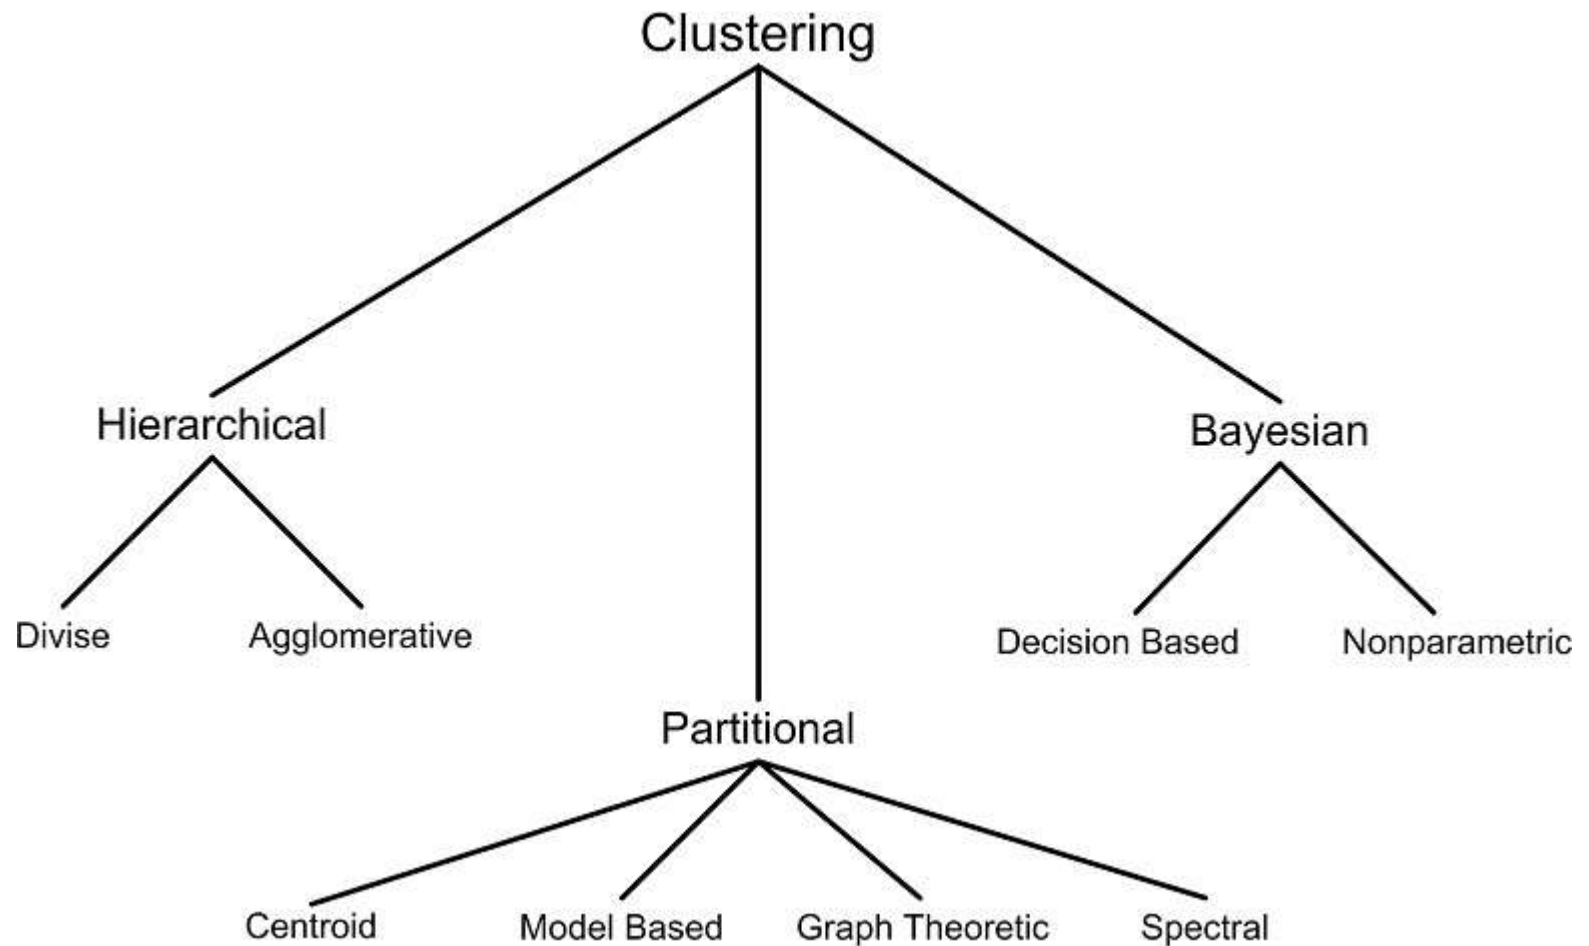
- **Intra-cluster cohesion** (compactness):
 - Cohesion measures how near the data points in a cluster are to the cluster centroid.
 - Sum of squared error (SSE) is a commonly used measure.
 - **Inter-cluster separation** (isolation):
 - Separation means that different cluster centroids should be far away from one another.
-

How many clusters?



- Possible approaches
 1. fix the number of clusters to k
 2. find the best clustering according to the criterion function (number of clusters may vary)

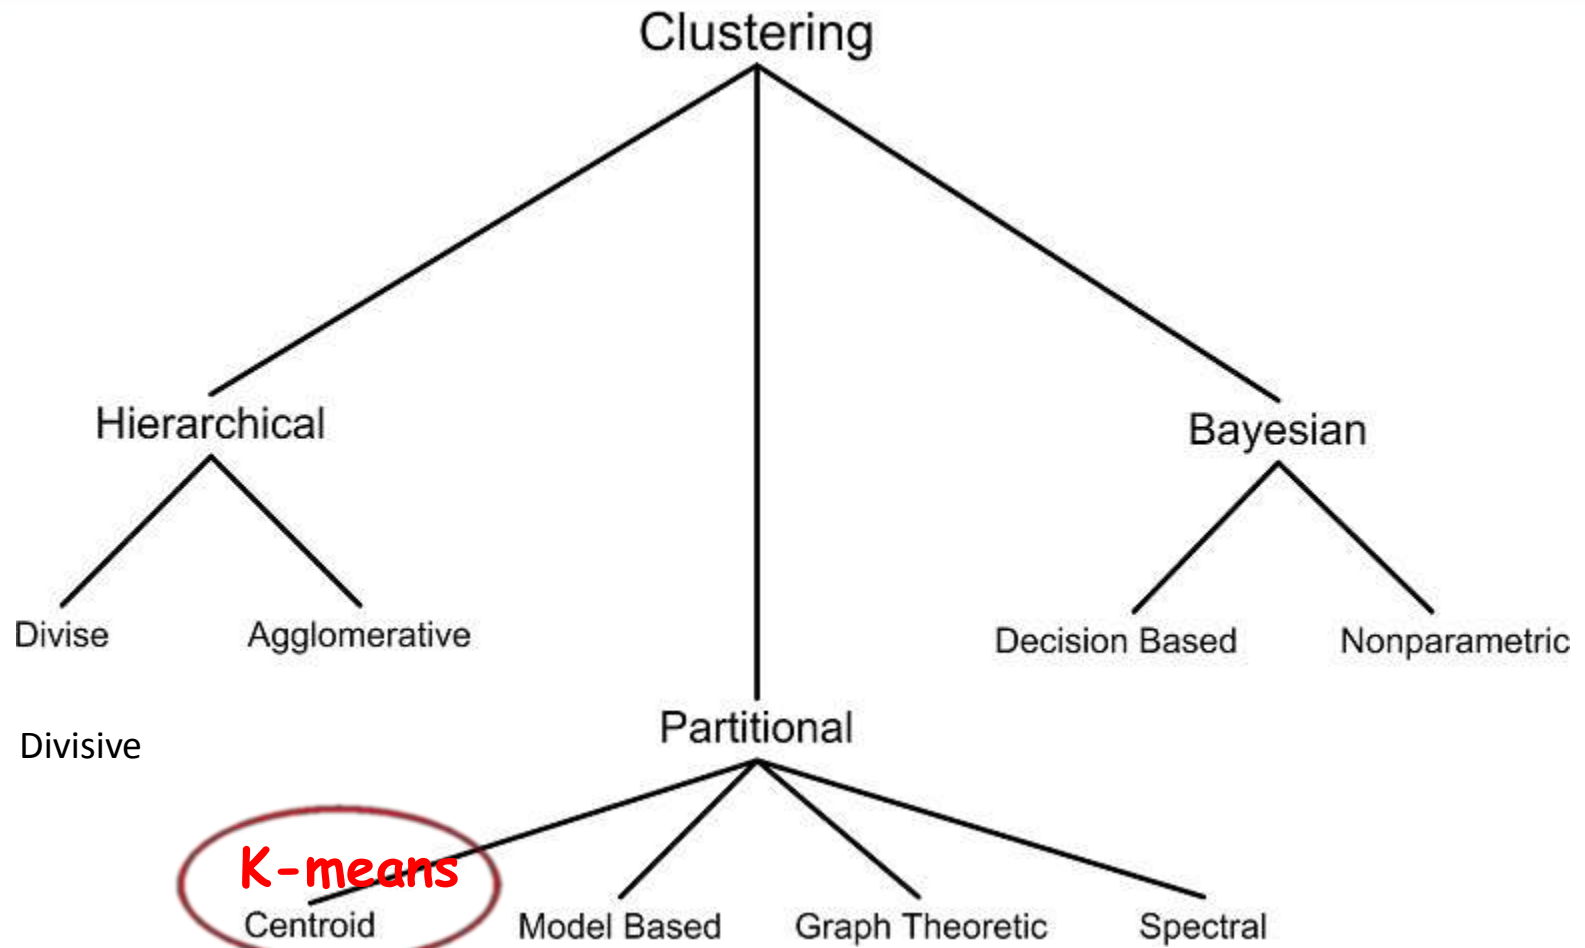
Clustering Techniques



Clustering Techniques

- **Hierarchical** algorithms find successive clusters using previously established clusters. These algorithms can be either **agglomerative** (“*bottom-up*”) or **divisive** (“*top-down*”):
 - ① **Agglomerative algorithms** begin with each element as a separate cluster and merge them into successively larger clusters;
 - ② **Divisive algorithms** begin with the whole set and proceed to divide it into successively smaller clusters.
 - **Partitional** algorithms typically determine all clusters at once, but can also be used as divisive algorithms in the hierarchical clustering.
 - **Bayesian** algorithms try to generate a *posteriori distribution* over the collection of all partitions of the data.
-

Clustering Techniques



K-Means clustering

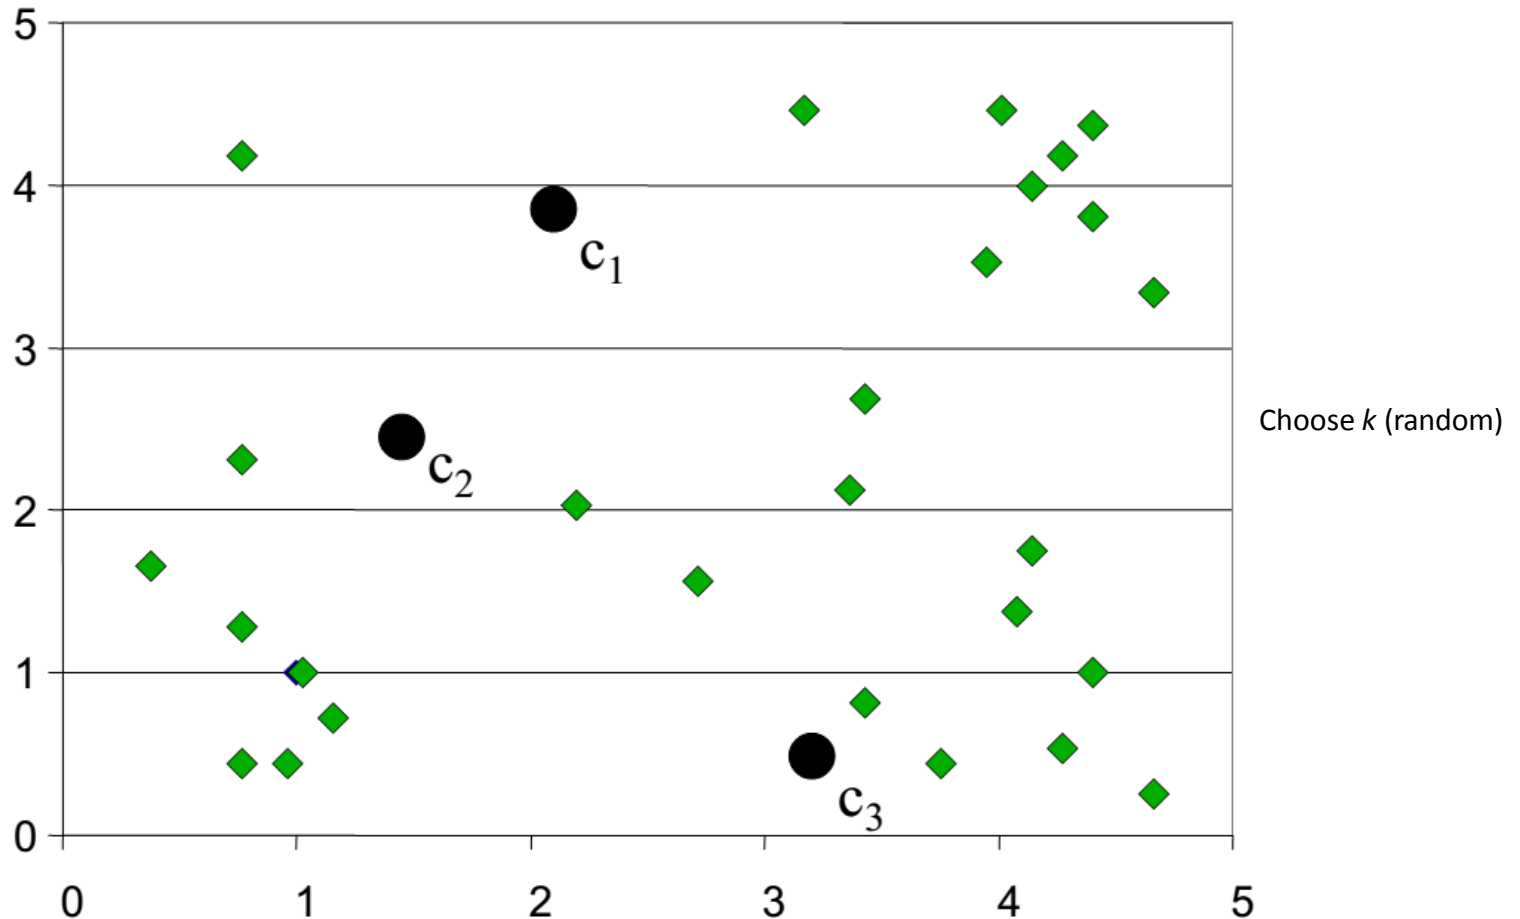
- K-means (MacQueen, 1967) is a **partitional clustering** algorithm
 - The k -means algorithm partitions the given data into k clusters:
 - Each cluster has a cluster **center**, called **centroid**.
 - k is specified by the user
-

K-means algorithm

- Given k , the *k-means* algorithm works as follows:
 1. Choose k (random) data points (**seeds**) to be the initial **centroids**, cluster centers
 2. Assign each data point to the closest **centroid**
 3. Re-compute the **centroids** using the current cluster memberships
 4. If a convergence criterion is not met, repeat steps 2 and 3
-

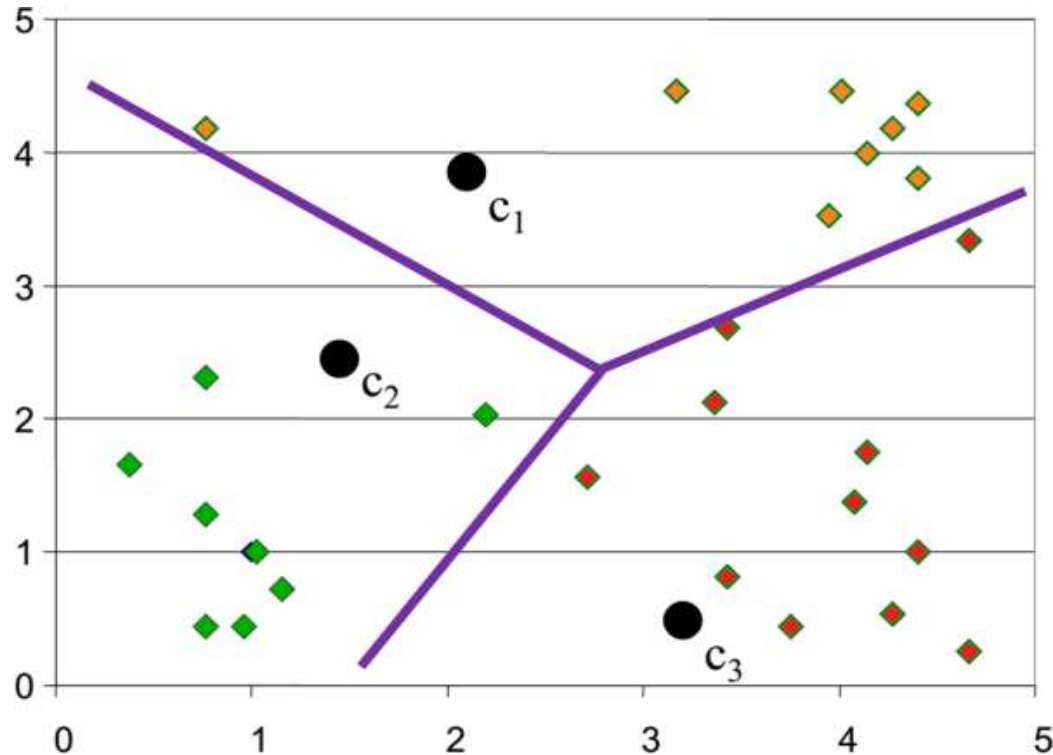
K-means clustering example: step 1

Randomly initialize the cluster centers (synaptic weights)



K-means clustering example – step 2

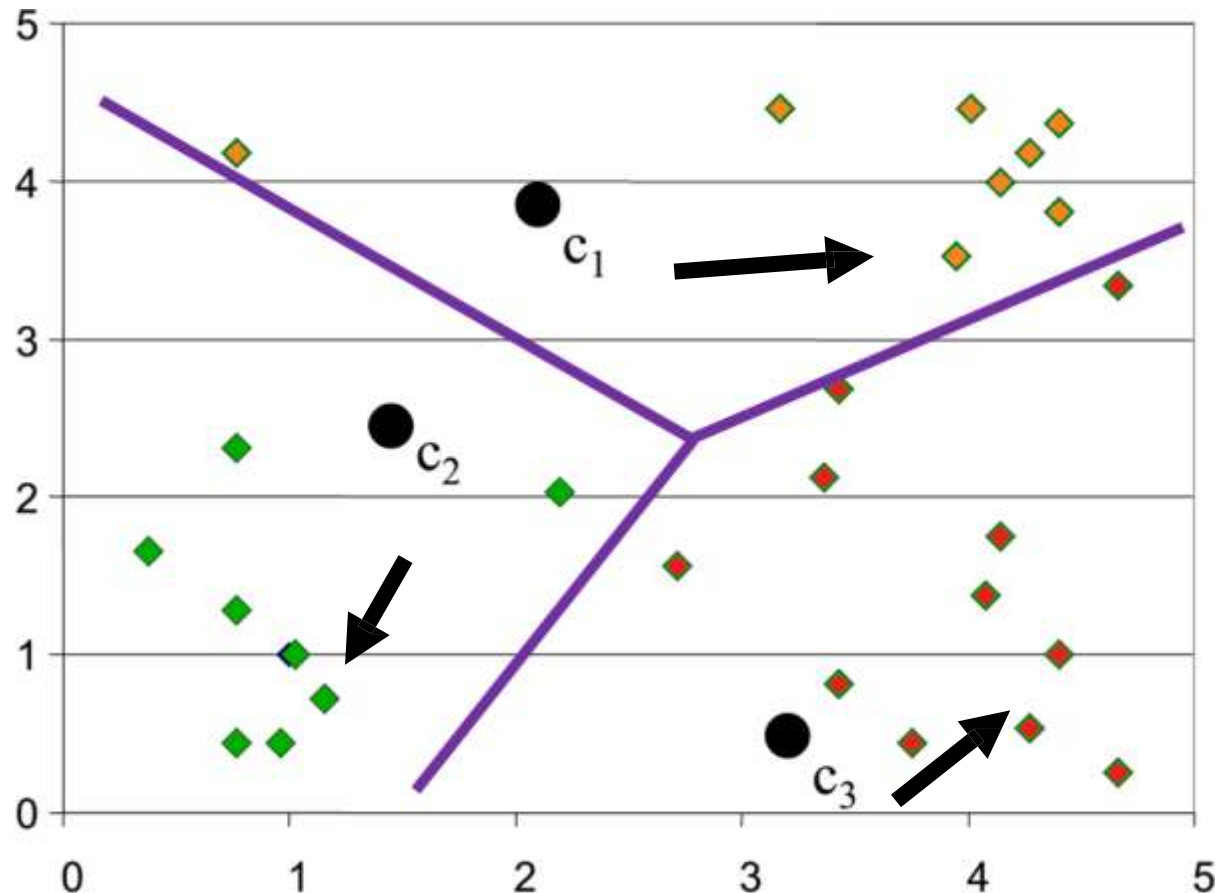
Determine cluster membership for each input
("winner-takes-all" inhibitory circuit)



Assign each data point to the closest **centroid**

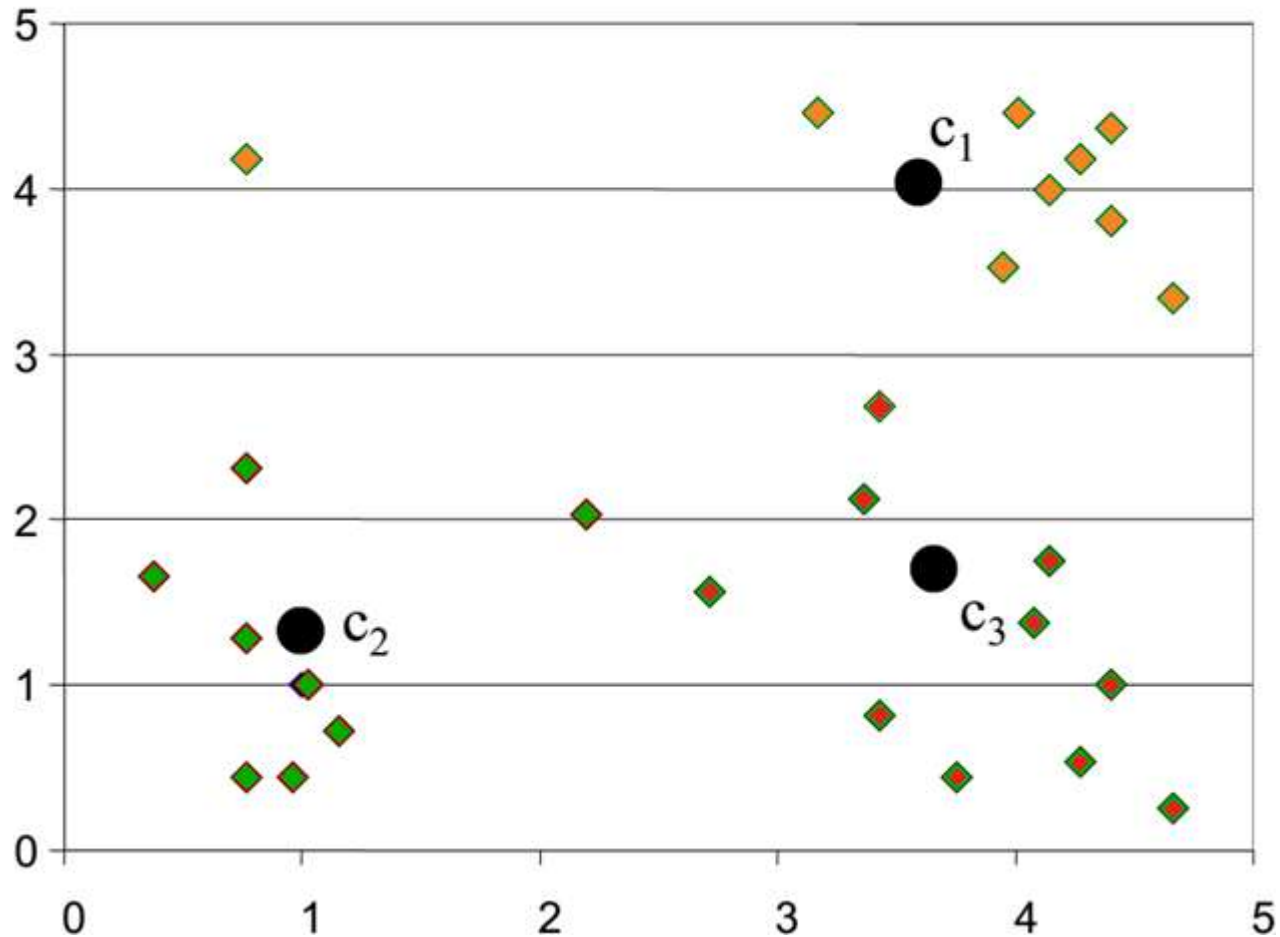
K-means clustering example – step 3

Re-estimate cluster centers (adapt synaptic weights)



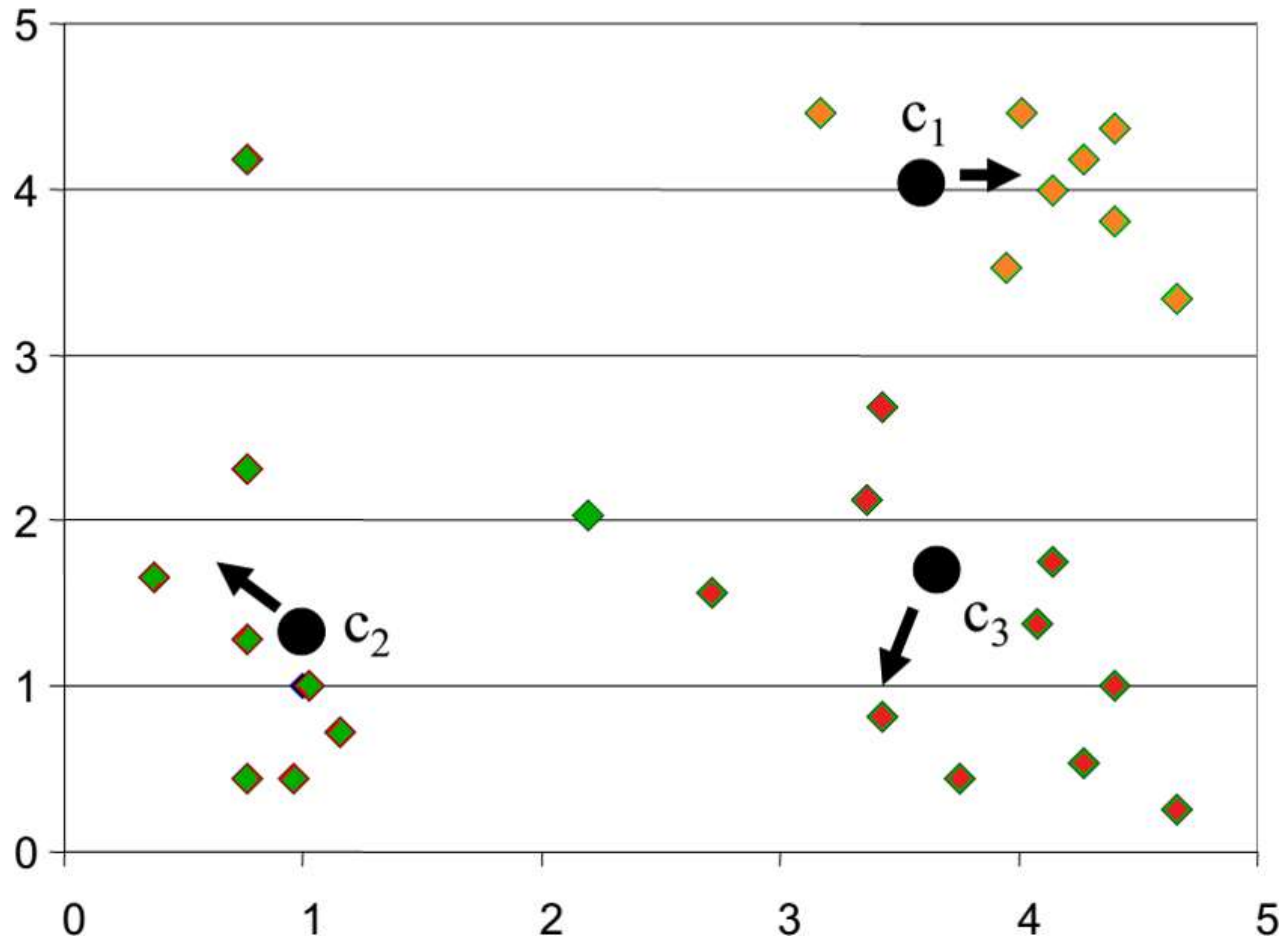
K-means clustering example

Result of first iteration



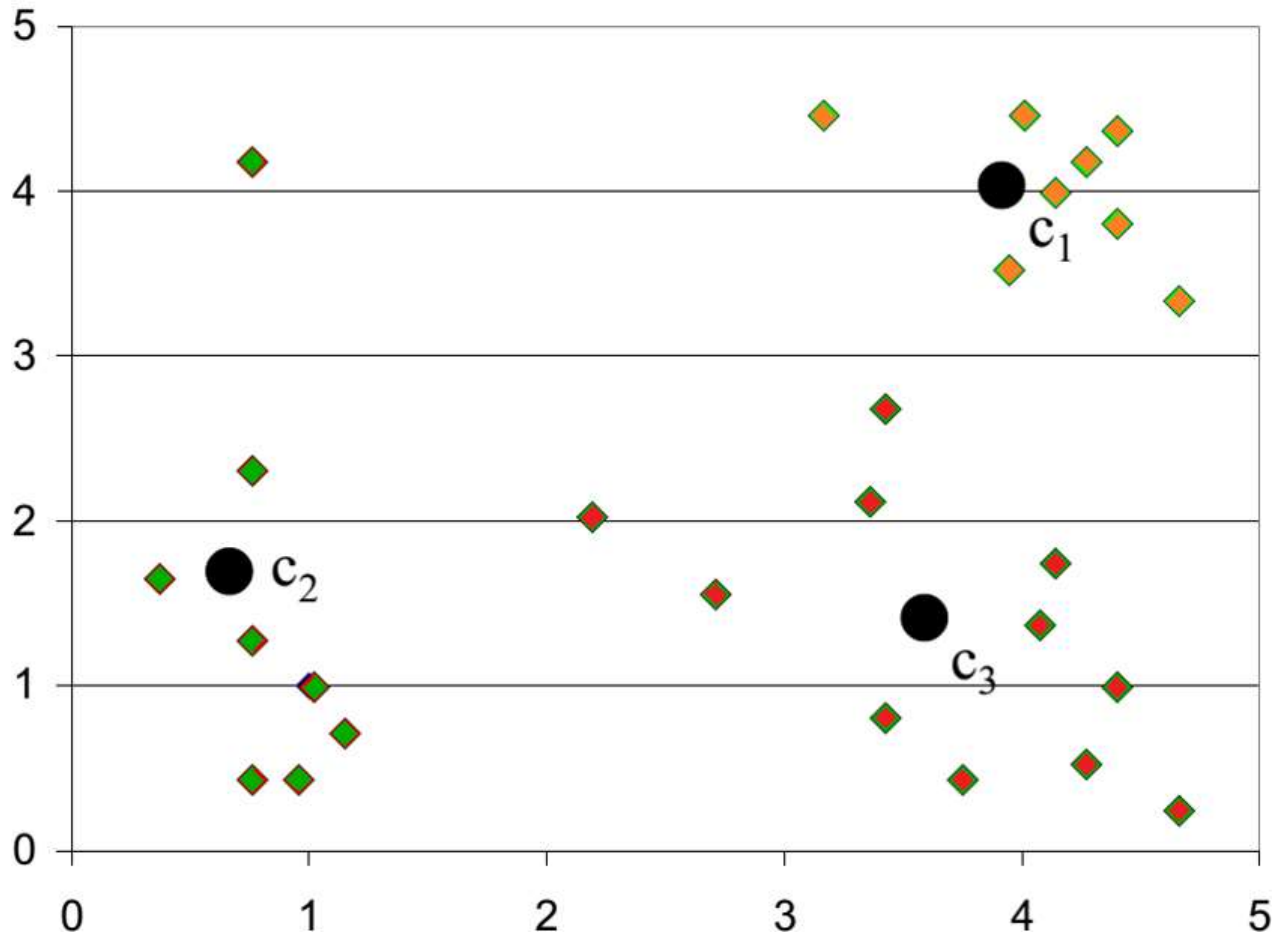
K-means clustering example

Second iteration



K-means clustering example

Result of second iteration



Why use K-means?

- Strengths:
 - Simple: easy to understand and to implement
 - Efficient: Time complexity: $O(tkn)$,
 - where n is the number of data points,
 - k is the number of clusters, and
 - t is the number of iterations.
 - Since both k and t are small. k -means is considered a linear algorithm.
 - K-means is the most popular clustering algorithm.
 - Note that: it terminates at a **local optimum** if SSE is used. The **global optimum** is hard to find due to complexity.
-

Weaknesses of K-means

- The algorithm is only applicable if the **mean** is defined.
 - For categorical data, *k*-mode - the centroid is represented by most frequent values.
 - The user needs to specify ***k***.
 - The algorithm is sensitive to **outliers**
 - Outliers are data points that are very far away from other data points.
 - Outliers could be errors in the data recording or some special data points with very different values.
-

K-means summary

- Despite weaknesses, *k*-means is still the most popular algorithm due to its simplicity and efficiency
 - No clear evidence that any other clustering algorithm performs better in general
 - Comparing different clustering algorithms is a difficult task. No one knows the correct clusters!
-

`Thank You!'
